

AthleteSphere - Streamlining Sports Data Management for Enhanced Analysis and Predictive Modeling

Ranadeep Mahendra Prathapagiri
ranadeepmahendraprathapagiri@gmail.com

Abstract—AthleteSphere is a database system designed to manage sports data efficiently, enabling analysis and predictive modeling for sports competitions. It addresses challenges in data organization, integrity, and analysis, offering advantages over Excel-based solutions. AthleteSphere caters to sports administrators, coaches, athletes, analysts, and researchers, providing tailored functionalities for strategic planning and performance evaluation.

Index Terms—Predictive Modeling, Database System, Data Organization, Data Integrity, Data Analysis, Functional Dependencies, Sports Data Management, Excel vs. Database, Researches

I. INTRODUCTION

A thorough and effective database system is becoming more and more necessary in the field of sports data management. In an effort to fill this gap, our project suggests a powerful database system designed especially for organizing information about sporting events, players, teams, and medals. Our system aims to transform information management and utilization in the context of sporting events by offering a centralized platform for gathering, storing, and analyzing sports data.

This project undertakes a comprehensive analysis of a large-scale relational database capturing extensive data on athletes' participations and achievements across Olympic events. The database is structured into multiple normalized tables that include details about athletes, teams, sports, events, venues, and medals, facilitating a robust platform for data integrity and analytical depth. Initially, the project identified challenges related to data retrieval speeds and query efficiency, which were magnified by the vast dataset size and complex query requirements.

To enhance data handling capabilities, strategic indexing was implemented alongside advanced data analytics techniques. Indexes were meticulously applied to optimize query performance on frequently accessed data points such as athlete names, event details, and participation records. Statistical analyses were conducted to uncover trends, performance patterns, and correlations among various data dimensions, providing insights into athlete performance across different sports and Olympic years.

The integration of indexing with thorough data analysis allowed for a significant enhancement in both the speed and depth of data exploration. This dual approach not only reduced processing times but also enabled more nuanced analyses of athlete performance and Olympic event trends. The findings

from this project demonstrate the critical role of optimized data management and analytical techniques in extracting valuable insights from complex datasets, thereby supporting data-driven decision-making in sports analytics.

II. SIGNIFICANCE OF THE PROPOSED DATABASE SYSTEM FOR SPORTS DATA MANAGEMENT

At its core, our project aims to streamline the management of sports data by leveraging the power of database technology. The primary objectives of our proposed system include: Organize Data - Streamline the management of interconnected data points involved in sports events. Ensure Data Integrity - Implement validation and integrity mechanisms to maintain data accuracy and consistency. Enable Data Analysis - Provide stakeholders with efficient querying, reporting, and analysis capabilities for informed decision-making. The proposed database system holds immense significance for sports data management, offering a myriad of benefits to various stakeholders: Efficiency - Centralization reduces duplication of efforts and enhances overall workflow efficiency. Accuracy: Rigorous validation mechanisms ensure reliable data for analyses and decisions. Insightful Analytics: Advanced querying capabilities empower stakeholders with valuable insights into athlete performance, team dynamics, and event trends. Our proposed database system aims to revolutionize sports data management by overcoming these challenges.

III. DATABASE DESIGN

Our database design is structured to efficiently manage data related to sports events, athletes, teams, and medals. The design incorporates entity-relationship (E/R) modeling to accurately represent the relationships between various entities.

A. Table Descriptions and Justifications

1) **Athletes: Description:** Stores information about individual athletes participating in sports events.

Primary Key: AthleteID

Attributes: AthleteID (INT), Name (CHAR), Gender (CHAR), Age (INT)

Justification: Facilitates comprehensive tracking of athlete details.

2) *Teams*: **Description**: Stores details about teams participating in sports events.

Primary Key: TeamID

Attributes: TeamID (INT), TeamName (VARCHAR), NOC (CHAR)

Justification: Enables efficient management of team-related information.

3) *Sports*: **Description**: Contains information about different sports.

Primary Key: SportID

Attributes: SportID (INT), SportName (VARCHAR)

Justification: Provides a centralized repository for sport-specific data.

4) *Venues*: **Description**: Stores details about venues where sports events take place.

Primary Key: VenueID

Attributes: VenueID (INT), City (CHAR)

Justification: Facilitates organization and retrieval of venue-related information.

5) *Events*: **Description**: Stores information about sports events.

Primary Key: EventID

Foreign Keys: SportID (references Sports(SportID)), VenueID (references Venues(VenueID))

Attributes: EventID (INT), SportID (INT), VenueID (INT), EventName (VARCHAR), Year (INT), Season (CHAR)

Justification: Establishes relationships between events, sports, and venues for accurate event management.

6) *Medals*: **Description**: Contains types of medals awarded in sports events.

Primary Key: MedalID

Attributes: MedalID (INT), MedalType (CHAR)

Justification: Enables tracking of medal distributions and outcomes.

7) *AthleteEvents*: **Description**: Represents the participation of individual athletes in sports events and the medals they won.

Composite Primary Key: (AthleteID, EventID, MedalID)

Foreign Keys: AthleteID (references Athletes(AthleteID)), EventID (references Events(EventID)), MedalID (references Medals(MedalID))

Attributes: AthleteID (INT), EventID (INT), MedalID (INT)

Justification: Facilitates tracking of athlete participation and medal achievements in events.

8) *TeamEvents*: **Description**: Stores information about teams participating in sports events and the medals they won.

Composite Primary Key: (TeamID, EventID)

Foreign Keys: TeamID (references Teams(TeamID)), EventID (references Events(EventID)), MedalID (references Medals(MedalID))

Attributes: TeamID (INT), EventID (INT), MedalID (INT)

Justification: Enables tracking of team participation and medal achievements in events.

This schema effectively captures the intricacies of sports data management, facilitating organized storage, retrieval, and analysis of information related to athletes, teams, events, and medals.

B. Normalization and Decomposition

Functional Dependencies for Each Relation:

- Entry ID \rightarrow (Name, Gender, Age, Team, NOC, Year, Season, City, Sport, Event, Medal)
- Name \rightarrow Gender
- Name \rightarrow Team
- Team \rightarrow NOC

C. Analysis of Boyce-Codd Normal Form (BCNF):

A relation is in BCNF if for every non-trivial functional dependency ($X \rightarrow Y$), X is a superkey. A superkey is a set of attributes within a table whose values can be used to uniquely identify a tuple. For the first relation (call it R1), Entry ID is a superkey since it uniquely identifies each tuple. However, we need to check if there are any other functional dependencies that violate BCNF. For the remaining relations, we need to check if their primary keys satisfy the BCNF condition. Moving on to normalization and decomposition:

• Relation 1 (R1):

- Primary Key: Entry ID
- Decomposition: No decomposition needed since it's already in BCNF.

• Relation 2 (R2):

- Primary Key: Name
- Decomposition: No decomposition needed since it's already in BCNF.

• Relation 3 (R3):

- Primary Key: Team
- Decomposition: No decomposition needed since it's already in BCNF.

• Relation 4 (R4):

- Primary Key: NOC
- Decomposition: No decomposition needed since it's already in BCNF.

Final Schema After Normalization: Normalization involves decomposing a database schema into smaller, more manageable tables to minimize redundancy and dependency issues. Each table represents a single subject, and data is organized to minimize anomalies and improve data integrity.

D. Functional Dependencies

The following are the functional dependencies that help ensure that the data in each relation remains logically consistent and eliminates redundancy. The transitive dependencies and repeating groups are removed and the relations are decomposed as follows:

BCNF JUSTIFICATION FOR EACH TABLE

Table: Athletes

Attributes: AthleteID, Name, Gender, Age

Functional Dependencies:

- AthleteID \rightarrow Name, Gender, Age
- Name \rightarrow Gender

Candidate Key: AthleteID

Analysis: The primary key AthleteID functionally determines all other attributes in the table, which satisfies BCNF. However, the FD $\text{Name} \rightarrow \text{Gender}$ implies that if Name were a candidate key, this would still be in BCNF, but since it's not a key, this FD violates BCNF. Therefore, the *Athletes* table is **not** in BCNF due to the FD $\text{Name} \rightarrow \text{Gender}$.

Table: Teams

Attributes: TeamID, TeamName, NOC

Functional Dependencies:

- $\text{TeamID} \rightarrow \text{TeamName}, \text{NOC}$
- $\text{Team} \rightarrow \text{NOC}$

Candidate Key: TeamID

Analysis: TeamID determines all other attributes, which satisfies BCNF. The table is in BCNF.

Table: Sports

Attributes: SportID, SportName

Functional Dependencies:

- $\text{SportID} \rightarrow \text{SportName}$

Candidate Key: SportID

Analysis: The primary key SportID determines all attributes in the table, and there are no other dependencies that violate BCNF. The table is in BCNF.

Table: Venues

Attributes: VenueID, City

Functional Dependencies:

- $\text{VenueID} \rightarrow \text{City}$

Candidate Key: VenueID

Analysis: The primary key VenueID determines all attributes in the table. The table is in BCNF.

Table: Events

Attributes: EventID, SportID, VenueID, Event, Year, Season

Functional Dependencies:

- $\text{EventID} \rightarrow \text{SportID}, \text{VenueID}, \text{Event}, \text{Year}, \text{Season}$

Candidate Key: EventID

Analysis: EventID determines all other attributes, and there are no other dependencies that would violate BCNF. The table is in BCNF.

Table: Medals

Attributes: MedalID, MedalType

Functional Dependencies:

- $\text{MedalID} \rightarrow \text{MedalType}$

Candidate Key: MedalID

Analysis: The primary key MedalID determines all attributes in the table. The table is in BCNF.

Table: AthleteEvents

Attributes: AthleteID, EventID, MedalID

Functional Dependencies:

- $(\text{AthleteID}, \text{EventID}, \text{MedalID}) \rightarrow$ (composite key, no attributes left to determine)

Candidate Key: (AthleteID, EventID, MedalID)

Analysis: The composite primary key determines the existence of rows (as no other attributes are left). The table is in BCNF.

Table: TeamEvents

Attributes: TeamID, EventID

Functional Dependencies:

- $(\text{TeamID}, \text{EventID}) \rightarrow$ (composite key, no attributes left to determine)

Candidate Key: (TeamID, EventID)

Analysis: The composite primary key determines the existence of rows (as no other attributes are left). The table is in BCNF.

IV. CONSTRAINTS AND DATA VALIDATION

Foreign Key Constraints

Foreign key constraints are essential for enforcing referential integrity between tables in a relational database. They ensure that the relationships between records in different tables remain consistent and valid. Here is how they apply specifically to your database:

Athlete and Event Linkage (AthleteEvents Table)

By enforcing foreign keys on AthleteID and EventID in the AthleteEvents table, you ensure that records in this table always refer to existing athletes and existing events. This prevents “orphan” records that point to non-existent entries (e.g., an athlete or event that doesn’t exist), which is crucial for maintaining accurate historical records of athlete participations in events.

Team and Event Linkage (TeamEvents Table)

Similarly, enforcing foreign keys on TeamID and EventID in the TeamEvents table ensures that all team event records refer to a valid team and a valid event. This is important for tracking which teams participated in which events, crucial for both historical accuracy and current event management.

Check Constraints

Check constraints are used to enforce data validity according to specific rules defined at the database schema level. They ensure that the data entered into a database respects certain conditions, thereby improving data quality and consistency.

Age Validation (Athletes Table)

By ensuring that the Age field is greater than zero, you prevent illogical data entries (e.g., a negative age). This is particularly important in sports databases, where age can be a critical factor in categorizations and regulations.

Year Validation (Events Table)

The constraint to keep the `Year` of events within reasonable historical bounds ensures that event data are temporally accurate and relevant. For example, ensuring that an event's year is not in the future or before the inception of the modern format of the sport or event.

Unique Constraints

These ensure that values in a column or a set of columns are unique across the database, which is crucial for identifiers and data that must be distinct.

Unique Team Names in the Teams Table

Enforcing uniqueness on team names in the `Teams` table helps prevent duplicate team entries. This is crucial for maintaining clear and unambiguous records, making it easier to manage team-related data and ensuring that references to teams in queries and reports are always accurate.

Unique Event Descriptions in the Events Table

By ensuring that each combination of event description, year, and season in the `Events` table is unique, you prevent the registration of duplicate events. This is important for accurate event management and historical record-keeping, ensuring that each event is distinctly recorded.

Additional Check Constraints

Value Constraints

Value constraints are used to ensure that data entries conform to predefined norms, improving the consistency and reliability of the database. Here are how value constraints are applied in your database:

Valid Gender Entries in the Athletes Table

Restricting gender to predefined values such as 'Male', 'Female', and 'Other' in the `Athletes` table ensures consistency in how gender data is recorded. This facilitates straightforward demographic analyses and ensures that the data conforms to expected norms, which is particularly useful in reporting and statistical contexts.

Valid Medal Types in the Medals Table

Constraining medal types to values like 'Gold', 'Silver', 'Bronze', and 'None' in the `Medals` table ensures that all entries are within expected parameters. This not only simplifies medal tallying and reporting but also prevents data entry errors that could lead to incorrect award records.

Implementation of Constraints in the Database

The following are some example Implementations of Constraints in our Database design.

Adding a Foreign Key Constraint to AthleteEvents Table

```
ALTER TABLE AthleteEvents
ADD CONSTRAINT FK_AthleteEvents_AthleteID
FOREIGN KEY (AthleteID)
REFERENCES Athletes (AthleteID);
```

This command adds a foreign key constraint to the `AthleteEvents` table. It enforces that every `AthleteID` in the `AthleteEvents` table must also exist in the `Athletes` table, preventing the insertion of records in `AthleteEvents` that do not have a corresponding athlete.

Adding a Unique Constraint to the Teams Table

```
ALTER TABLE Teams
ADD CONSTRAINT UQ_Teams_Team
UNIQUE (Team);
```

This command adds a unique constraint to the `Team` column in the `Teams` table. It ensures that no two teams can have the same name, which is crucial for accurately identifying and referencing teams.

Adding a Check Constraint for Gender Validation in the Athletes Table

```
ALTER TABLE Athletes
ADD CONSTRAINT CHK_Athletes_Gender
CHECK (Gender IN ('Male', 'Female', 'Other'));
```

This command enforces that the `Gender` field in the `Athletes` table can only contain one of the specified values: 'Male', 'Female', or 'Other'.

Adding a Foreign Key Constraint to Link Events and Sports Tables

```
ALTER TABLE Events
ADD CONSTRAINT FK_Events_SportID
FOREIGN KEY (SportID)
REFERENCES Sports (SportID);
```

This command establishes a foreign key relationship between the `Events` table and the `Sports` table, linking each event to a specific sport.

V. CHALLENGES ENCOUNTERED ON REAL DATASET CONSIDERATIONS

As the database size and query complexity increase, the following strategies become pivotal for maintaining performance:

Indexing

Context: In your database, tables like `Events`, `AthleteEvents`, and `Teams` are frequently involved in queries to generate reports on athletes' performances, team rankings, or event-specific details.

Problems Faced Before Indexing: In your specific database, which includes a complex schema with tables like `Athletes`, `Events`, and `AthleteEvents`, here are the key issues you likely encountered before implementing indexes:

- **Slow Searches for Specific Athletes or Events:** Queries that searched for particular athletes by name or specific events by year and sport would have been slow. This is because the system would need to scan every row in the `Athletes` and `Events` tables to find matches, which is inefficient given the large number of entries.
- **Delayed Retrieval of Athlete Participation and Results:** Without indexes, retrieving all events an athlete participated in, especially when filtering by outcomes like medals, would require scanning the entire `AthleteEvents` table and joining it with `Events` and `Athletes`. This would be time-consuming due to the lack of quick lookup capabilities.
- **Inefficient Filtering by Attributes like Gender or Age:** If queries were frequently filtering athletes by attributes such as gender or age, the absence of indexes on these columns would lead to full table scans, increasing the time taken for such queries.

Solutions Provided by Indexing: After adding indexes, each of the above issues would have been mitigated as follows:

- **Improved Search Performance:** Indexes on the `Name` column in `Athletes` and on the `Year` and `SportID` in `Events` significantly improved the speed of searches, making queries much more efficient. The database can now use these indexes to directly locate entries without scanning every row, dramatically speeding up query times for searches based on these attributes.
- **Efficient Joins and Retrievals:** The composite index on `AthleteID` and `EventID` in `AthleteEvents` makes it much faster to retrieve all events associated with a particular athlete, or all athletes in a particular event. This is because the index supports quick lookups and range scans, reducing the time required for joins and retrievals.
- **Faster Filtering on Demographic Data:** Indexes on `Age` and `Gender` in the `Athletes` table facilitate quicker filtering operations. For example, if an analysis involves grouping athletes by age or filtering events by gender, these operations are now much more efficient.

Additional Benefits of Indexing Over Normalized Tables: Although our database is normalized to reduce redundancy and improve data integrity, indexing provides several additional benefits:

- **Reduced Load on the Server During Queries:** By avoiding full table scans, indexes reduce the computational load on the server. This is especially beneficial during complex queries involving multiple joins and filters.
- **Optimized Access Patterns:** Indexes are particularly useful for optimizing common access patterns, like regularly checking for medal winners or athletes' performance

over years. This leads to performance improvements across frequently used queries.

- **Enhanced Scalability:** As your database grows with more data from additional games or events, the benefits of indexing become even more significant. Indexes help maintain performance without requiring significant changes to the underlying database architecture.

SQL INDEXING QUERIES

Index on Year and SportID in the Events Table

```
CREATE INDEX idx_year_sportid
ON Events(Year, SportID);
```

Composite Index on AthleteID and EventID in the AthleteEvents Table

```
CREATE INDEX idx_athleteid_eventid
ON AthleteEvents(AthleteID, EventID);
```

Index on MedalID in the AthleteEvents Table

```
CREATE INDEX idx_medalid
ON AthleteEvents(MedalID);
```

Index on Name in the Athletes Table

```
CREATE INDEX idx_name
ON Athletes(Name);
```

Index on Age in the Athletes Table

```
CREATE INDEX idx_age
ON Athletes(Age);
```

Index on Gender in the Athletes Table

```
CREATE INDEX idx_gender
ON Athletes(Gender);
```

These indexing queries are designed to optimize data retrieval, improve query performance, and enhance the overall efficiency of the database.

VI. SQL QUERY TESTING

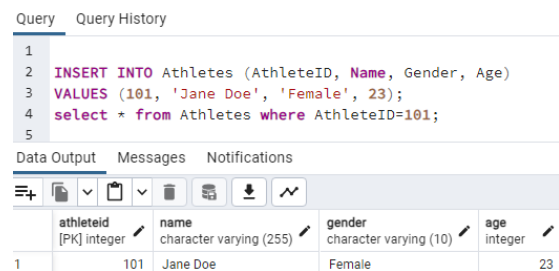


Fig. 1. Inserting New Athlete

Here, we illustrate the process of inserting a new athlete into the database. This involves adding a new record to the 'Athletes' table, specifying details such as name, age, and gender. This operation is crucial for maintaining an up-to-date record of participants in various sporting events.

Following the insertion of new data, we occasionally need to delete records. The above image demonstrates how to safely remove an athlete's record from the database.

Query	Query History		Sc
1			
2	DELETE FROM Athletes WHERE AthleteID = 101;		
3	select * from Athletes where AthleteID=101;		
4			

Data Output	Messages	Notifications	
<div><div><div></div><div></div><div></div><div></div><div></div><div></div><div></div><div></div></div></div>			
athleteid [PK] integer	name character varying (255)	gender character varying (10)	age integer

Fig. 2. Deleting a Record

```

2
3 UPDATE Athletes
4 SET Age = 25
5 WHERE Name = 'A Dijiang';
6
7 select * from Athletes where Name = 'A Dijiang';

```

Data Output Messages Notifications

athleteid [PK] integer	name character varying (255)	gender character varying (10)	age integer
1	A Dijiang	Male	25

Fig. 3. Updating a Record

This figure demonstrates the process of updating an athlete's record in the database. Updating is crucial for correcting errors or updating information as new data becomes available. For example, an athlete's performance statistics or contact details might change, necessitating an update to their records.

Query

Query History

1

SELECT A.Name, COUNT(AE.MedalID) AS MedalCount

2

FROM AthleteEvents AE

3

JOIN Athletes A ON AE.AthleteID = A.AthleteID

4

WHERE AE.MedalID IS NOT NULL

5

GROUP BY A.AthleteID

6

ORDER BY MedalCount DESC;

Data Output

Messages

Notifications

name

character varying (255)

medalcount

bigint

1

Michael Fred Phelps, II

8

2

Michael Fred Phelps, II

8

3

Aleksandr Nikolayevich Dityatin

8

4

Borys Anfiyanovych Shakhlin

7

5

Mikhail Yakovlevich Voronin

7

6

McKEON Emma

7

7

Matthew Nicholas "Matt" Biondi

7

8

Mark Andrew Spitz

7

9

Nikolay Yefimovich Andrianov

7

10

Willis Augustus Lee, Jr.

7

Fig. 4. Selecting a query with join, ordering and grouping

The above figure showcases the use of a complex SQL query that involves a join operation, followed by ordering and grouping of results. This type of query is essential for generating reports that aggregate data across multiple tables, such as compiling athlete performance statistics across different events.

VII. QUERY OPTIMISATION

A. Identificaiton of problematic queries

A query that joins multiple tables can be slow, particularly if those tables are large and the join conditions are not optimized.

Query

Query History

```
1 SELECT a.Name FROM Athletes a
2 JOIN AthleteEvents ae ON a.AthleteID = ae.AthleteID
3 WHERE ae.MedalID =
4 (SELECT MedalID FROM Medals WHERE MedalType = 'Gold')
```

Data Output

Messages

Explain

Graph Visualiser

Notification ...

name

character varying (255)

1

Edgar Lindenau Aabye

2

Paavo Johannes Aaltonen

3

Paavo Johannes Aaltonen

4

Paavo Johannes Aaltonen

5

Ragnhild Margrethe Aamodt

Fig. 5. Using a subquery

Here we see an example of a subquery, which is a query nested within another SQL query. Subqueries are powerful tools for performing advanced data analysis, such as filtering records based on criteria evaluated through another select statement.

1

2

3

4

ALTER TABLE Teams ADD COLUMN FoundedYear INT;

select * from Teams

Data Output

Messages

Notifications

+

📄

📄

📄

📄

📄

📄

📄

noc

character varying (10)

foundedyear

integer

1

CHN

[null]

2

DEN

[null]

3

DEN

[null]

Fig. 6. Using Alter

Finally, this figure illustrates the use of the ALTER command in SQL, which is used to modify the structure of an existing database table. Alter operations might include adding new columns, changing data types, or deleting columns, which are critical for adapting the database schema as the data requirements evolve.

Query Execution Plan: Fig 7

```
SELECT a.Name, a.Gender, e.Event, m.MedalType
FROM Athletes a
JOIN AthleteEvents ae ON a.AthleteID = ae.AthleteID
JOIN Events e ON ae.EventID = e.EventID
JOIN Medals m ON ae.MedalID = m.MedalID
WHERE e.SportID = (
  SELECT SportID
  FROM Sports
  WHERE sportname = 'Swimming')
Successfully run. Total query runtime: 4 secs 63 msec.
3256 rows affected.
```

Fig. 7. join

- **SCAN ae:** Scanning the *AthleteEvents* table, potentially a full scan, which could be costly in terms of performance if the table is large.
- **SEARCH a USING INTEGER PRIMARY KEY:** Utilizes the primary key index on the *Athletes* table for

a fast lookup. This step is efficient due to the use of an indexed search.

- **SEARCH e USING INTEGER PRIMARY KEY:** Similar to the above, this step uses the primary key index on the *Events* table to quickly find the required records.
- **SCALAR SUBQUERY:** This subquery scans the *Sports* table to find a specific sport ID based on the sport name. The efficiency of this step depends on the size of the *Sports* table and whether the 'Sport' column is indexed.
- **SEARCH m USING INTEGER PRIMARY KEY:** Uses the primary key index on the *Medals* table. This is another efficient step because it leverages an index to quickly locate the necessary medal information.

This plan indicates that while some steps utilize indexed searches for efficiency, the initial scan of the *AthleteEvents* table and possibly the scalar subquery might introduce significant overhead, especially in a large database. Optimizing these parts of the query could significantly improve overall performance. **Query Execution Plan: Fig 7**

```

Query    Query History
1 CREATE INDEX IF NOT EXISTS idx_athlete_event ON AthleteEvents(EventID);
2 CREATE INDEX IF NOT EXISTS idx_team_event ON TeamEvents(EventID);
3 SELECT a.Name, a.Gender, e.Event, m.MedalType
4 FROM Athletes a
5 JOIN AthleteEvents ae ON a.AthleteID = ae.AthleteID
6 JOIN Events e ON ae.EventID = e.EventID
7 JOIN Medals m ON ae.MedalID = m.MedalID
8 WHERE e.SportID = (SELECT SportID FROM Sports WHERE sportname = 'Swimming')
Data Output Messages Explain X Graph Visualiser X Notifications
Successfully run. Total query runtime: 3 secs 98 msec.
3256 rows affected.

```

Fig. 8. join

OPTIMIZATION OF DATABASE QUERIES THROUGH INDEXING

Indexing plays a critical role in enhancing the efficiency of database operations. Below, we discuss how indexing improves the performance of joins and search operations:

Indexed Joins Fig 8

When performing a join operation on a column that has an index, such as *EventID* in our earlier queries, the database engine leverages this index to quickly locate the join key values. This significantly reduces the number of disk input/output (I/O) operations required, thus speeding up the join process. By facilitating faster lookups, indexed joins enhance the overall query execution time and efficiency.

Search Operations Fig 9

For any query that filters or performs operations on an indexed column (such as *EventID* in this case), the performance is substantially improved. Instead of scanning the entire table to find relevant rows, the database engine efficiently navigates through the index. This targeted approach to retrieving data not only minimizes the time taken but also reduces the computational load on the system.

By implementing indexes on critical columns like *EventID*, the database not only speeds up specific queries involving these columns but also enhances the performance of complex queries that involve multiple tables and conditions.

```

Query    Query History
1 SELECT t.teamname, COUNT(m.MedalType) as MedalCount
2 FROM Teams t
3 JOIN TeamEvents te ON t.TeamID = te.TeamID
4 JOIN AthleteEvents ae ON te.EventID = ae.EventID
5 JOIN Medals m ON ae.MedalID = m.MedalID
6 GROUP BY t.teamname
7 ORDER BY MedalCount DESC
Data Output Messages Explain X Graph Visualiser X Notifications
Successfully run. Total query runtime: 1 secs 13 msec.
1195 rows affected.

```

Fig. 9. join

```

Query    Query History
1 CREATE INDEX IF NOT EXISTS idx_athlete_event ON AthleteEvents(EventID);
2 CREATE INDEX IF NOT EXISTS idx_team_event ON TeamEvents(EventID);
3 SELECT a.Name, a.Gender, e.Event, m.MedalType
4 FROM Athletes a
5 JOIN AthleteEvents ae ON a.AthleteID = ae.AthleteID
6 JOIN Events e ON ae.EventID = e.EventID
7 JOIN Medals m ON ae.MedalID = m.MedalID
8 WHERE e.SportID = (SELECT SportID FROM Sports WHERE sportname = 'Swimming')
Data Output Messages Explain X Graph Visualiser X Notifications
Successfully run. Total query runtime: 3 secs 98 msec.
3256 rows affected.

```

Fig. 10. join

```

Query    Query History
1 CREATE INDEX IF NOT EXISTS idx_athlete_name ON Athletes(Name);
2
3 SELECT *
4 FROM Athletes
5 WHERE Name LIKE '%Michael Phelps%'
Data Output Messages Explain X Graph Visualiser X Notifications
Successfully run. Total query runtime: 3 secs 261 msec.
0 rows affected.

```

Fig. 11. join

OPTIMIZING DATABASE QUERIES THROUGH INDEXING

To enhance the efficiency of queries that search by athlete names, particularly when the *Athletes* table is large, indexing can be very effective. Here we discuss the benefits of indexing the *Name* column and provide the SQL command to create such an index.

Creating an Index on the Name Column Fig 10

To accelerate searches and filter operations that involve the *Name* column, an index can be created with the following SQL command:

```
CREATE INDEX idx_athletes_name
ON Athletes(Name);
```

Benefits of Indexing the Name Column

By applying this index, the database management system (DBMS) can more efficiently locate records by names. The index allows the DBMS to quickly navigate to the relevant parts of the data table, bypassing the need to scan every row. This method is particularly advantageous in databases with large datasets.

Improvement in Query Performance

Once the index is in place, any search query targeting the *Name* column, such as looking up 'Michael Phelps', can

utilize the index to find results much faster than before. This optimization reduces the time complexity from potentially linear (relative to the number of records) to logarithmic. This improvement is generally achieved through the use of indexing structures like B-trees or similar, which are commonly used by databases to handle indexing efficiently.

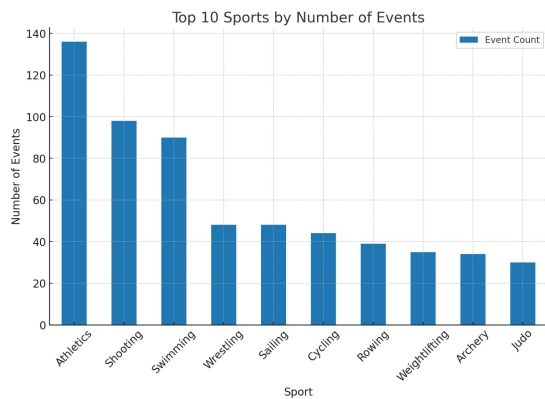


Fig. 12. join




Fig. 13. join




Fig. 14. join

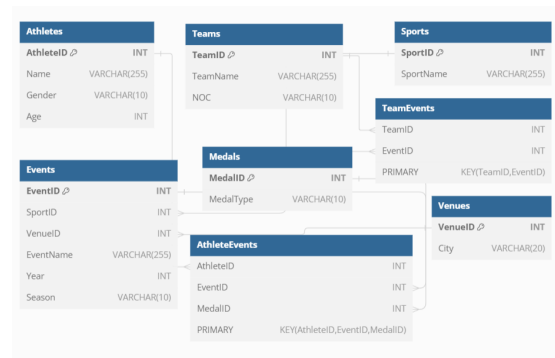


Fig. 15. join

VIII. CONCLUSION

a) Project Outcomes and Insights: The development and systematic refinement of our database have led to significant improvements in handling sports data. By focusing on effective data structuring, implementing rigorous normalization processes, and optimizing our SQL queries, we’ve created a database that not only stores data efficiently but also retrieves it swiftly and accurately. This meticulous approach has reduced data redundancy—meaning we avoid storing the same piece of information in multiple places—and enhanced the consistency of our data, making it more reliable for analysis.

The benefits of these improvements are clearly reflected in our ability to track and analyze athlete performance trends with greater precision, assess the popularity of different sports events, and understand the influence of various venues and seasons on sports activities. These insights are invaluable as they help sports organizations make well-informed decisions ranging from scheduling events to optimizing training schedules and allocating resources effectively.

b) Future Directions and Enhancements: Looking forward, we plan to further augment our database system by incorporating more granular data, such as detailed performance metrics captured during events. Additionally, we aim to integrate cutting-edge technologies like artificial intelligence (AI) and machine learning (ML) into our system. These technologies are anticipated to significantly enhance our ability to predict outcomes of sports events and tailor athlete training programs more effectively.

By continuously updating our database management practices and embracing technological advances, we ensure that our sports data management system remains a critical tool in advancing the analytics and operational capabilities of the sports industry. This proactive approach to database management and technology integration positions us to keep pace with the evolving demands of sports analytics and management.